Deep Learning Performance Checklist

Data Preparation
- ☐ Store as TFRecords
- ☐ Reduce size of input data
- ☐ Use TensorFlow Datasets

Data Reading
- ☐ Use tf.data
- ☐ Prefetch data
- ☐ Parallelize CPU processing
- ☐ Parallelize I/O and processing
- ☐ Enable nondeterministic ordering
- ☐ Cache data
- ☐ Turn on experimental optimizations
- ☐ Autotune parameter values

Data Augmentation
- ☐ Use GPU for augmentation

Training
- ☐ Use automatic mixed precision
- ☐ Use larger batch size
- ☐ Use multiples of eight
- ☐ Find the optimal learning rate
- ☐ Use `tf.function`
- ☐ Overtrain, then generalize
    - ☐ Progressive sampling
    - ☐ Progressive augmentation
    - ☐ Progressive resizing
- ☐ Install an optimized stack for the hardware
- ☐ Optimize number of parallel CPU threads
- ☐ Use better hardware
- ☐ Distribute training
- ☐ Examine industry benchmarks

Inference
- ☐ Use an efficient model
- ☐ Quantize the model
- ☐ Prune the model
- ☐ Use fused operations
- ☐ Enable GPU persistence